

Rag Concepts

Retrieval-Augmented Generation combines retrieval systems with large language models. Documents are first loaded and split into chunks. Embedding models convert chunks into vectors. Vector databases such as Chroma, FAISS, and Qdrant store embeddings. Retrievers search for relevant chunks using similarity search. Retrieved context is passed to the language model to generate grounded responses. Good chunking and retrieval quality are critical for strong RAG systems.